



Intro to Apache Spark

XQuery Working Group

Text Mining at Scale

Fall 2019

What is it?



Apache Spark is an open-source distributed general-purpose cluster-computing framework

What does that
even mean?



Break it down

- open-source (non-proprietary, works well with others)
- distributed (more than one computer, often remote, entails communication costs)
- general-purpose (can do any big data operation)
- cluster-computing (handle parallel programming)

MapReduce Model

- Grew out of needs of Google to process big data quickly
- Across a distributed parallel computer cluster
- TL; DR
 - Input data
 - Map (filtering and sorting operations)
 - Reduce (summary operations like counting)
 - Output data

MapReduce Model: Input

- Start with a set of key/value pairs
- Where have we seen key/value pairs?
 - JSON
 - Need to convert XML to JSON (or JSONL)

MapReduce Model: Map

Map the input pairs to a function and output intermediate pairs

```
map (input_key, input_value) -> list(output_key, intermediate_value)
```

Function typically a sorting or filtering function

Output an intermediate file

MapReduce Model: Shuffle

- Hidden phase where the intermediate output from the Map function is transferred to the Reduce function
- Apache Spark optimize the shuffle part, so you only have to focus on the map and reduce parts
- Apache Spark handles managing task scheduling and fault tolerance

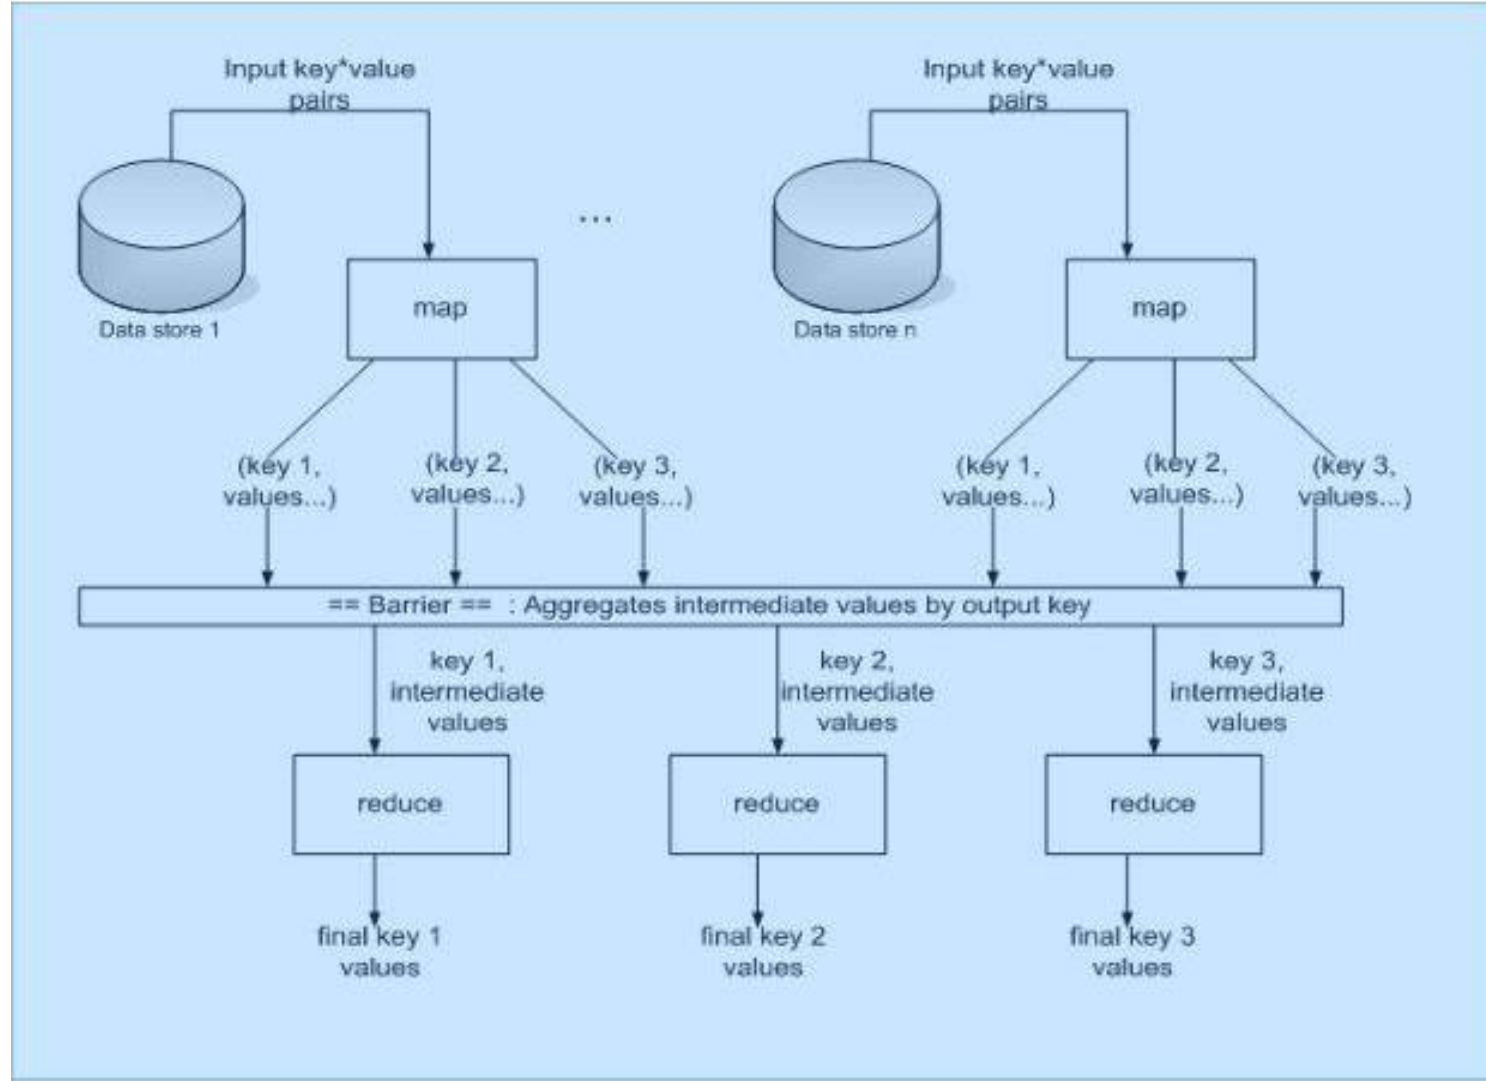
MapReduce Model: Reduce

Reduce the intermediate pairs to the output key/value pair

`reduce (output_key, list(intermediate_value)) -> list(output_value)`

Collect the intermediate outputs and perform a summary function on them up to produce a single result

Summary functions include sum, count, max, min, avg



Trade-offs

- Trade off between computation and communication costs
 - Communication costs usually more expensive than computational costs
- ∴ Spark isn't always the fastest, but handles big data the best

Pseudo-Code Example: Word Count

Map_function (String name, String document):

for each word w **in** document:

 return (w, 1)

Reduce function (word, partialCounts):

 sum = 0

for each pc **in** partialCounts:

 sum += pc

 return (word, sum)

RumbleDB

- Can query big data on the Spark cluster
- Contained in JSONL files
- Using JSONiq query language
- And familiar FLWOR expressions
- <http://rumbledb.org/>

Language Game

<https://namethatlanguage.org/ntl>

Language Game Dataset

- 16 million records of each guess, one JSON record per line:

```
{ "target": "Turkish",  
  "sample": "af0e25c7637fb0dc56fac6d49aa55e",  
  "choices": [  
    "Hindi",  
    "Lao",  
    "Maltese",  
    "Turkish"  
  ],  
  "guess": "Maltese",  
  "date": "2013-08-19",  
  "country": "AU"  
}
```